8. Fully-Connected Neural Networks

First, we will employ the most basic form of a deep neural network, in which the neurons in adjacent layers are fully connected to one another.

You will be working in the files`part2-mnist/nnet_fc.py`**in this problem**

Training and Testing Accuracy Over Time

0.0/1.0 punto (calificable)

We have provided a toy example **nnet_fc.py** in which we have implemented for you a simple neural network. This network has one hidden layer of 10 neurons with a rectified linear unit (ReLU) nonlinearity, as well as an output layer of 10 neurons (one for each digit class). Finally, a softmax function normalizes the activations of the output neurons so that they specify a probability distribution. Reference the PyTorch Documentation and read through it in order to gain a better understanding of the code. Then, try running the code on your computer with the command `python3 nnet_fc.py`. This will train the network with 10 epochs, where an epoch is a complete pass through the training dataset. Total training time of your network should take no more than a couple of minutes. At the end of training, your model should have an accuracy of more than $\%85$ on test data.

Note: We are not using a softmax layer because it is already present in the loss: PyTorch's `nn.CrossEntropyLoss` combines `nn.LogSoftMax` with `nn.NLLLoss`.

Report the test accuracy below.

Test Accuracy =

Enviar Ha realizado 0 de 3 intentos

Improving Accuracy

0.0/5.0 puntos (calificable)

We would like to try to improve the performance of the model by performing a mini grid search over hyper parameters (note that a full grid search should include more values and combinations). To this end, we will use our **baseline model (batch size 32, hidden size 10, learning rate 0.1, momentum 0 and the ReLU activation function)** and modify one parameter each time while keeping all others to the baseline. We will use the validation accuracy of the model after training for 10 epochs. For the LeakyReLU activation function, use the default parameters from pyTorch (negative_slope=0.01).

Note: If you run the model multiple times from the same script, make sure to initialize the **numpy and pytorch random seeds to 12321 before each run**.

Which of the following modifications achieved the highest validation accuracy?

- baseline (no modifications)
- batch size 64
- learning rate 0.01
- momentum 0.9
- LeakyReLU activation

What is the validation accuracy of the best performing model?

Validation Accuracy =

Does the model variation that achieved the highest validation accuracy achieved also the highest test accuracy?

- Yes
- No

Enviar Ha realizado 0 de 3 intentos

Improving Accuracy - Hidden 128

0.0/3.0 puntos (calificable)

Modifying the model's architecture is also worth considering. Increase the hidden representation size from 10 to 128 and repeat the grid search over the hyper parameters. This time, what modification achieved the highest validation accuracy?

- baseline (no modifications)
- batch size 64
- learning rate 0.01
- momentum 0.9
- LeakyReLU activation

Enviar Ha realizado 0 de 3 intentos

Discusión

Ocultar Discusión

Tema: Unit 3 Neural networks (2.5 weeks):Project 3: Digit recognition (Part 2) / 8. Fully-Connected Neural Networks

Add a Post

Mostrar todas las publicaciones por actividad reciente

[Staff]can you grant an extension, I was travelling heavliy this week Dear Staff, please, grant a day of extension. I am half away through this project, it would be ni...	3
Hidden 128 -- possible Grader issue? I got **Improving Accuracy** part correct. And I presume to get the correct answer for **Imp...	10
[STAFF] May I get extension for the Midterm Hi, I'm travelling with my family from Aug-2nd to Aug-10th, by any chance, may I get extensio...	3
STAFF： HELP me check with my result I got correct with the first two questions but incorrect with the improving accuracy (hidden 12...	24
nnet_cnn.py ValueError I'm unable to run the code without getting the following error in Training_utilities line 31, in c...	3
[Staff] Running nnet_fc.py returns error	3
[staff] help with torch installation.. I installed torch after many tries. Then it started giving trouble with the numpy. I uninstalled n...	4
mnist_model_fully_connected.pt model file Hi. Given following line in train_model: > torch.save(model, 'mnist_model_fully_connected.pt')...	4
Small differences between the validation and the test, my be data related, what should we learn from this experiment ? Small differences between the validation and the test, my be data related, what should we lea...	7
Hidden 128 I'm new to PyTorch, and having read the portions of the documentation that seem relevant, I'...	8

Will extension be granted for this project? Dear Staff, could you please grant an extension for this project like the previous ones. Thank ...		2
Hidden 128 For the last question and using the latest version of Anaconda and torch under Windows 10, I...		1
Don't use GPU for grading purpose Just a heads up, using GPU with PyTorch will yield you slightly different training and testing re... Community TA		6
For the LeakyReLU activation function, use the default parameters from pyTorch (negative_slope=0.01). How to change this parameter? In which part of the code should this be?		2
Accuracy value In question "improving accuracy" I got all the questions right apart from actual value of accur...		2
[STAFF] Training and Testing Accuracy Over Time - Can't reach more than 85% on test set I have created a network as requested: Number of inputs 784, Number of hidden units 10, N...		6
Grid Search. Hello. When doing grid search we have to find the best combination or we return everything t...		3
Why do we call the validation data as dev_data? I notice that in all the code 10% of data is used for validation and is extracted from the MNIST...		2
[staff] nnet_fc.py running very slowly Did anyone else's nnet_fc.py run very slowly? I didn't change any of the parameters on mine a...		4
Any specific reason for 784 in nn.Linear(784, 10) ? 784 is the number of inputs of training examples, right? Is there a reason we chose specificall...		2

Load more